

Off-Whitepaper

Ethereum

Micah Dameron

*Beautiful is better than ugly.
Explicit is better than implicit.
Simple is better than complex.
Complex is better than complicated.*

The Zen of Python

Abstract

The goal of this paper^a is to create and expand concepts from Ethereum about which, notwithstanding any earlier documentation, there may be some justified confusion. We use pseudocode rather than mathematical notation to describe Ethereum's operation, because pseudocode has many advantages when describing ABSTRACT STATE MACHINES,^b like Ethereum. This paper takes an approach to describing Ethereum that focuses on clarity and approachability. Our prime source has been the Ethereum *Yellowpaper*, but much supplemental knowledge has been found elsewhere and crucial points from other sources have been added as well for the reader's benefit.

^aFormally, *Blanched-Almond Paper*

^bE. Borger and S. Robert F., *Abstract state machines: A method for high-level system design and analysis*. 1, pp. 3-8. Springer, 2003.

Acknowledgements

Thank you to the Ethereum founders for creating a product worth writing about in minute detail. Thanks to the Ethereum Foundation for maintaining this product in its basic integrity. Thanks to the ConsenSys Mesh for supporting my work on this project, by contributing your vast knowledge and expertise. Finally, thank you to Dr. Gavin Wood for your technical astuteness and creative genius; your Yellowpaper has given Ethereum a soul worth decoding.

Contents

| | |
|--|----------|
| 1. Native Currency & Mining | 5 |
| 1.1. Mining | 5 |
| 1.1.1. Total Difficulty | 5 |
| 1.2. EVM Code | 5 |
| 1.3. Programming Languages | 5 |
| 1.3.1. Lower-Level Lisp | 5 |
| 1.3.2. Solidity | 5 |
| 1.3.3. Serpent | 5 |
| 1.3.4. Viper | 5 |
| 2. Memory and Storage | 5 |
| 2.1. Data Structures | 5 |
| 2.1.1. World State | 5 |
| 2.2. The Block | 5 |
| 2.2.1. Block Header | 6 |
| 2.2.2. Block Footer | 6 |
| 2.3. State Database | 6 |
| 2.3.1. RLP | 6 |
| 2.3.2. Account State | 7 |
| 2.4. Bloom Filter | 7 |
| 2.4.1. Transaction Receipts | 7 |
| 3. Processing and Computation | 7 |
| 3.1. State Transition Function | 7 |
| 3.1.1. Ethash | 8 |
| 3.2. Verification | 8 |
| 3.2.1. Ommers | 8 |
| 3.3. Transactions | 8 |
| 3.4. Execution | 8 |
| 3.4.1. Intrinsic Validity | 8 |
| 3.4.2. Execution Model | 9 |
| 3.5. Substate | 9 |
| 3.5.1. Message Calls | 9 |
| 3.5.2. Contract Creation | 10 |
| 3.5.3. Account Creation | 10 |
| 3.6. Halting | 10 |
| 3.7. Gas | 10 |
| 3.7.1. Machine State | 11 |
| 3.7.2. Exceptional Halting | 11 |
| 3.7.3. EVM Code | 11 |
| 3.7.4. Opcodes/EVM Assembly | 12 |

| | |
|-------------|----|
| I. Appendix | 13 |
| A. Opcodes | 13 |
| References | 15 |
| Glossary | 16 |
| Acronyms | 17 |

1. Native Currency & Mining

Because Ethereum is turing-complete, there needs to be a *network price unit* to mitigate the possibility of abusing the network with excessive computational expenditures. The smallest unit of currency in Ethereum is the Wei, which is equal to $\Xi 10^{-18}$. All currency transactions in Ethereum are counted in Wei. There is also the Szabo, which is $\Xi 10^{-6}$, and the Finney, which is $\Xi 10^{-3}$.

| Unit | Ether | Wei |
|--------|----------------------------|---------------------------|
| Ether | $\Xi 1.000000000000000000$ | 1,000,000,000,000,000,000 |
| Finney | $\Xi 0.001000000000000000$ | 1,000,000,000,000,000 |
| Szabo | $\Xi 0.000001000000000000$ | 1,000,000,000,000 |
| Wei | $\Xi 0.000000000000000001$ | 1 |

1.1. Mining

1.1.1. Total Difficulty

The *Total Difficulty* of a block is defined recursively by a function which calculates the difficulty of all blocks prior to the header in the present block.

| Pseudocode | Definition |
|---|---|
| <code>presentstate(total.difficulty)</code> | Total difficulty of <i>this block</i> . |
| <code>presentstate(block.parent)</code> | <i>This block's</i> parent block. |
| <code>presentstate(block.difficulty)</code> | <i>This block's</i> difficulty. |

1.2. EVM Code

1.3. Programming Languages

1.3.1. Lower-Level Lisp

The Lisp-Like low level language: a human-writable language used for authoring simple contracts and trans-compiling to higher-level languages.

1.3.2. Solidity

A JavaScript-like language, and the most commonly used language for creating smart contracts in Ethereum.

1.3.3. Serpent

1.3.4. Viper

2. Memory and Storage

2.1. Data Structures

Merkle-Patricia Trees Merkle-Patricia Trees

2.1.1. World State

Also known simply as "state", this is a MAPPING of **addresses** and **account states** (RLP data structures), this is also known as *state*, or σ . This mapping is not stored on the blockchain, rather it is stored as a Merkle-Patricia **trie** in a DATABASE BACKEND^a that maintains a mapping of bytearrays to bytearrays.^b The cryptographic internal data going back to the **root node** represents the *State* of the Blockchain at any given root, i.e. at any given *time*.^c As a whole, the state is the sum total of database relationships in the **state database**. The state is an inert position on the chain, a position between prior state and post state; a block's frame of reference, and a defined set of relationships to that frame of reference.

2.2. The Block

A block is made up of 17 different elements. The first 15 elements are part of what is called the *block header*.

2.2.1. Block Header

Notation : header

Description : The information contained in a block besides the transactions list. This consists of:

1. **Parent Hash** – This is the Keccak-256 hash of the parent block's header.
2. **Ommers Hash** – This is the Keccak-256 hash of the ommer's list portion of this block.

^aA database backend is accessed by users indirectly through an external application, most likely an Ethereum client; see also: **state database**
^bA bytearray is specific set of bytes [data] that can be loaded into memory. It is a structure for storing binary data, e.g. the contents of a file.

^cThis permanent data structure makes it possible to easily recall any previous state with its root hash keeping the resources off-chain and minimizing on-chain storage needs.

3. **Beneficiary** – This is the 20-byte address to which all block rewards are transferred.
4. **State Root** – This is the Keccak-256 hash of the root node of the state trie, after a block and its transactions are finalized.
5. **Transactions Root** – This is the Keccak-256 hash of the root node of the trie structure populated with each transaction from a Block's transaction list.
6. **Receipts Root** – This is the Keccak-256 hash of the root node of the trie structure populated with the receipts of each transaction in the transactions list portion of the block.
7. **Logs Bloom** – This is the bloom filter composed from indexable information (log address and log topic) contained in the receipt for each transaction in the transactions list portion of a block.
8. **Difficulty** – This is the difficulty of this block – a quantity calculated from the previous block's difficulty and its timestamp.
9. **Number** – This is a quantity equal to the number of ancestor blocks behind the current block.
10. **Gas Limit** – This is a quantity equal to the current maximum gas expenditure per block.
11. **Gas Used** – This is a quantity equal to the total gas used in transactions in this block.
12. **Timestamp** – This is a record of Unix's time at this block's inception.
13. **Extra Data** – This byte-array of size 32 bytes or less contains extra data relevant to this block.
14. **Mix Hash** – This is a 32-byte hash that verifies a sufficient amount of computation has been done on this block.
15. **Nonce** – This is an 8-byte hash that verifies a sufficient amount of computation has been done on this block.
16. **Ommers Block Headers** – These are the same components listed above for any ommers.

2.2.2. Block Footer

17. **Transaction Series** – This is the only non-header content in the block.

2.3. State Database

2.3.1. RLP

Notation : rlp

Description : RLP encodes arrays of nested binary data to an arbitrary depth; it is the main serialization method for data in Ethereum. RLP encodes mainly structure and does not pay heed to what type of data it is encoding.

Positive RLP integers are represented with the most significant value stored at the lowest memory address (big endian) and without any leading zeroes. As a result, the RLP integer value for 0 is represented by an empty byte-array. If a non-empty deserialized integer begins with leading zeros it is invalid.⁷

The global state database is encoded as RLP for fast traversal and inspection of data. In structure it constitutes a mapping between *addresses* and *account states*. Since it is stored on node operator's computers, the tree can be traversed speedily and without network delay. RLP encodes values as byte-arrays, or as sequences of further values.¹

This means that:

```

if    rlp(x)           = bytearray
then  rlp(bytearray)   = true
elif  rlp(x)           = value
then  rlp(value)       = true
elif  rlp(x)           = null
then  rlp(x)           = false

```

1. If the RLP-serialized byte-array contains a single byte integer value less than 128, then the output is exactly equal to the input.

In other words:

2.3.2. Account State

Notation : body

Description : The EVM-code fragment that executes each time an account receives a message call.

Description : The account state is made up of four variables:

1. **nonce** The number of transactions sent from this address, or the number of contract creations made by the account associated with this address.
2. **balance** The number of Wei owned by this address.
3. **storage_root** A 256-bit (32-byte) hash of the root node of a Merkle Patricia tree that encodes the storage contents of the account.
4. The storage root aspect of an account's state is the hash of the trie
5. **code_hash** The hash of the EVM code of this account's contract.

The account state is the state of any particular account during some specified world state.

Nonce The **nonce** aspect of an ACCOUNT'S STATE is the number of transactions sent from, or the number of contract-creations by, the address of that account.^a

Storage Root The **storage root** aspect of an ACCOUNT'S STATE is the hash of the trie^b

Code Hash The **code hash** aspect of an ACCOUNT'S STATE is the HASH OF THE EVM CODE of this account. Code hashes are STORED in the **state database**. Code hashes are permanent and they are executed when the address belonging to that account RECEIVES a message call.

Balance The amount of **Wei OWNED** by this account.

- Key/value pair stored inside the root hash.
- L_I^* , is defined as the element-wise transformation of the base function
- The *element-wise transformation of the base-function* refers to all of the key/value pairs in L_I
- L_I refers to a particular **trie**.

^a σ is the world state at a certain given time, and n is the number of transactions or contract creations by that account.

^bA particular path from root to leaf in the **state database** that encodes the STORAGE CONTENTS of the account.

2.4. Bloom Filter

Notation : logs_bloom

Description : The Bloom Filter is composed from indexable information (logger address and log topics) contained in each log entry from the receipt of each transaction in the transactions list.

2.4.1. Transaction Receipts

3. Processing and Computation

3.1. State Transition Function

State Transitions come about through a what is known as the State Transition Function; this is an abstraction of several operations in Ethereum which comprise the overall act of computing changes to the *machine state* prior to adding them to the *world state*, that is, through them being finalized and rewards applied to a given miner. `apply_rewards` and `block_beneficiary` are here. subsectionMining

Block Beneficiary The 160-bit (20-byte, or 20-character) address to which all fees collected from the successful mining of a block are transferred.

Apply Rewards The third process in `block_finalization` that sends the mining reward to an account's address. A scalar value corresponding to the difficulty level of a current block. This can be calculated from the previous block's difficulty level and the timestamp.

3.1.1. Ethash

GHOST Protocol

3.2. Verification

Verifies Ommers headers

3.2.1. Ommers

Ommershash

3.3. Transactions

The basic method for Ethereum accounts to interact with each other. Transactions lie at the heart of Ethereum, and are entirely responsible for the dynamism and flexibility of the platform. Transactions are the bread and butter of state transitions, that is of block additions, which contain all of the computation performed in one block. Each transaction applies the execution changes to the *machine state*, a temporary state which consists of all the temporary changes in computation that must be made before a block is finalized and added to the world state.

Notation : sender

Description : A function that maps transactions to their sender using ECDSA of the SECP-256k1 curve, (excepting the latter three signature fields) as the datum to sign. The sender of a given transaction can be represented: `transaction.sender`

3.4. Execution

Description : The execution of a transaction defines the state transition function: `stf`. However, before any transaction can be executed it needs to go through the initial tests of intrinsic validity.

3.4.1. Intrinsic Validity

The criteria for intrinsic validity are as follows:

- The transaction follows the rules for *well-formed RLP* (recursive length prefix.)
- The *signature* on the transaction is valid.
- The *nonce* on the transaction is valid, i.e. it is equivalent to the sender account's current nonce.
- The *gas_limit* is greater than or equal to the *intrinsic_gas* used by the transaction.
- The sender's account balance contains the cost required in up-front payment.

Accordingly, the post-transactional state of Ethereum is expressed thus:

```
transaction(post.state) = stf(present.state,
```

^aDesignated "intrinsic_gas" in the Yellowpaper

```
transaction)
```

While the amount of gas used in the execution is expressed: `stf(gas_used)` and the accrued log items belonging to the transaction are expressed: `stf(logsbloom, content)(logsbloom, set)` Information concerning the result of a transaction's execution is stored in the transaction receipt `tx_receipt`. The set of log events which are created through the execution of the transaction, `logs_set` in addition to the bloom filter which contains the actual information from those log events `logs_bloom` are located in the transaction receipt. In addition, the post-transaction state `post_transaction(state)` and the amount of gas used in the block containing the transaction receipt `post(gas_used)` are stored in the transaction receipt. Thusly the transaction receipt is a record of any given execution.

A valid transaction execution begins with a permanent change to the state: the nonce of the sender account is increased by one and the balance is decreased by the *collateral_gas*^a which is the amount of gas a transaction is required to pay prior to its execution. The original transactor will differ from the sender if the message call or contract creation comes from a contract account executing code.

After a transaction is executed, there comes a PROVISIONAL STATE:

```
post_execution(provisional.state)
```

Gas used for the execution of individual EVM opcodes prior to their potential addition to the *world_state* creates the provisional state. *productive_gas*, and an associated substate *substate_a*.

Code execution always depletes gas. If gas runs out, an out-of-gas error is signaled (*oog*) and the resulting state defines itself as an empty set; it has no effect on the world state. This describes the transactional nature of Ethereum. In order to affect the *WORLD STATE*, a transaction must go through completely or not at all.

3.4.2. Execution Model

Description : The stack-based *virtual machine* which lies at the heart of the Ethereum and per-

forms the actions of a computer. This is actually an instantial runtime that executes several substates, as EVM computation instances, before adding the finished result, all calculations having been completed, to the final state via the finalization function.

In addition to the system state σ , and the remaining gas for computation g , there are several pieces of important information used in the execution environment that the execution agent must provide; these are contained in the tuple I :

- `account_address`, the address of the account which owns the code that is executing.
- `sender_address` the sender address of the transaction that originated this execution.
- `originator_price` the price of gas in the transaction that originated this execution.
- `input_data`, a byte array that is the input data to this execution; if the execution agent is a transaction, this would be the transaction data.
- `account_address` the address of the account which caused the code to be executing; if the execution agent is a transaction, this would be the transaction sender.
- `newstate_value` the value, in Wei, passed to this account if the execution agent is a transaction, this would be the transaction value.¹
- `code.array` the byte array that is the machine code to be executed.¹
- `samestate_header` the block header of the present block.
- `the stack depth` the depth of the present message-call or contract-creation (i.e. the number of CALLs or CREATEs being executed at present).¹

3.5. Substate

Description : A smaller, temporary state that is generated during transaction execution. It contains three sets of data:

- The `accounts` tagged for self-destruction following the transaction's completion.
`self_destruct(accounts)`

- The `logs_series`, which creates checkpoints in EVM code execution for frontend applications to explore, and is made up of `the_logs_set` and `logs_bloom` from the `tx_receipt`.
- The refund balance.^a

3.5.1. Message Calls

Description :

Notation : `message_call`

Description : A message call can come from a transaction or internally from contract code execution. It contains the field `DATA`, which consists of user input to a message call. Messages allow communication between accounts (whether contract or external,) and are a carryover from established concepts in Computer Science, most notably the *MPI: Message-Passing Framework*. Messages can come in the form of `msg_calls` which give output data. If an account has EVM code in it (a contract account,) this code gets executed when the account receives a message call. Message calls and contract creations are both *transactions*, but contract creations are never considered the same as message calls. Message calls always transfer some amount of value to an account. If the message call is an account creation transaction then the value given is taken on the role of an endowment toward the new account. Every time an account receives a message call it returns the body, something which is triggered by the `init` function. A message call can come through a transaction, or through the internal execution of code. Message call transactions only contain data. They are separate from regular, standard *transactions*.

Message calls always have a universally agreed-upon cost in gas. There is a strong distinction between contract creation transactions and message call transactions. Computation performed, whether it is a contract creation or a message call, represents the currently legal valid state. There can be no invalid transactions from this point.¹ There is also a message call/contract creation *stack*. This stack has a depth, depending on how many transactions are in

^aThe `SSTORE` operation increases the amount refunded by resetting contract storage to zero from some non-zero state.

it. Contract creations and message calls have entirely different ways of executing, and are entirely different in their roles in Ethereum. The concepts can be conflated. Message calls can result in computation that occurs in the next state rather than the current one. If an account that is currently executing receives a message call, no code will execute, because the account might exist but has no code in it yet. To execute a message call transactions are required:

- Sender
- Transaction_Originator
- Recipient
- Account (usually the same as the recipient)
- Available_Gas
- Value
- Gas_Price
- An arbitrary length byte-array. `arb_array`
- Present_Depth of the message call/contract creation stack.

Notation : `data`

Description : User data input to a `message_call`, structured as an unlimited size byte-array.

3.5.2. Contract Creation

Notation : `init`

Description : When `INIT` is executed it returns the `BODY`. `Init` is executed only once at `ACCOUNT_CREATION`, and permanently discarded after that. Contract creation transactions are equal the recursive length prefix of an empty byte-sequence.

3.5.3. Account Creation

3.6. Halting

Execution Environment

Notation : `ERE`

Description : The environment under which an Autonomous Object executes in the EVM: the EVM runs as a part of this environment.

Notation : `big_endian_f`

Description : `BIG_ENDIAN_FUNCTION` This function expands a positive-integer value to a big-endian byte array of minimal length. When accompanied by a `.` operator, it signals sequence concatenation. The `big_endian` function accompanies RLP serialization and deserialization.

3.7. Gas

Description : The fundamental network cost unit converted to and from Ether as needed to complete the transaction while it is sent. Gas is arbitrarily determined at the moment it is needed, by the block and according to the miners decision to charge certain fees.

Miner Choice Miners choose which gas prices they want to accept.

Gasprice

Notation : `gas_limit`

Description : A value equal to the current limit of gas expenditure per block, according to the miners.

Gaslimit Any unused gas is refunded to the user.

Gasused

Description : A value equal to the total gas used in transactions in this block.

3.7.1. Machine State

The machine state is a tuple consisting of five elements:

1. `gas_available`
2. `program_counter`
3. `memory_contents` A series of zeroes of size 2^{256}

4. `memory_words.count`

5. `stack_contents`

There is also, `[to_execute]`: the current operation to be executed

3.7.2. Exceptional Halting

An exceptional halt may be caused by a handful of boolean values:

```
forall instruction.x
if gas_empty = true
then signal halt
elif instruction.x = fake
then signal halt
elif stack = terse
then signal halt
elif jumpdest = bad
then signal halt
else exec instruction.x
```

```
forall instruction.y
[...]
[...]
[...]
[...]
```

```
forall instruction.z
[...]
[...]
[...]
[...]
```

then signal `controlled_halt`

No instruction can, through its execution, cause an exceptional halt. They can only happen if some instruction, for whatever reason, fails to execute.

- The amount of remaining gas in each transaction is extracted from information contained in the `machine_state`
- A simple iterative recursive loop¹ with a boolean value:

true indicating that in the run of computation, an exception was signaled

false indicating in the run of computation, exceptions were signaled. If this value remains false for the duration of the execution until the set of transactions becomes a series (rather than an empty set.) This means that the machine has reached a controlled halt.

Substate The substate is an emergent, ever-changing ball of computational energy that is about to be applied to the main state. It is the *meta state* by which transactions are decided valid and to be added to the blockchain.

3.7.3. EVM Code

The bytecode that the EVM can natively execute. Used to explicitly specify the meaning of a message to an account.

Notation : contract

Description : A piece of EVM Code that may be associated with an Account or an Autonomous Object.

3.7.4. Opcodes/EVM Assembly

The human readable version of EVM code. But what exactly are these computer instructions that can be executed with the same level of veracity and certainty as Bitcoin transactions? How do they come about, what makes them up, how are they kept in order, and what makes them execute? The first part of answering this question is understanding opcodes. In traditional machine architectures, you may not be introduced to working with processor-level assembly instructions for some time. In Ethereum however, they are essential to understanding the protocol because they are the most minute and subtle (yet HUGEY important) things going on in the Ethereum Blockchain at any moment, and they are the real "currency," that Ethereum trades in. I'll explain what I mean by that in a minute. First, let's go over a few Opcodes:^a

^aA full list of Opcodes is in Appendix B

The STOP Opcode is used in order to stop a computation once it has completed, or to halt a computation if it has run out of gas. The ADD, MUL, SUB, and DIV operations are addition, multiplication, subtraction and division operations. The In/Out columns refer to inputs (to `machine_state`), the state which decides every new `world_state`.

Part I.

Appendix

A. Opcodes

| Data | Opcode | Gas | Input | Output |
|------|--------------|-----|-------|--------|
| 0x00 | STOP | 0 | 0 | 0 |
| 0x01 | ADD | 3 | 2 | 1 |
| 0x02 | MUL | 5 | 2 | 1 |
| 0x03 | SUB | 3 | 2 | 1 |
| 0x04 | DIV | 5 | 2 | 1 |
| 0x05 | SDIV | 5 | 2 | 1 |
| 0x06 | MOD | 5 | 2 | 1 |
| 0x07 | SMOD | 5 | 2 | 1 |
| 0x08 | ADDMOD | 8 | 3 | 1 |
| 0x09 | MULMOD | 8 | 3 | 1 |
| 0x0a | EXP | 10 | 2 | 1 |
| 0x0b | SIGNEXTEND | 5 | 2 | 1 |
| 0x10 | LT | 3 | 2 | 1 |
| 0x11 | GT | 3 | 2 | 1 |
| 0x12 | SLT | 3 | 2 | 1 |
| 0x13 | SGT | 3 | 2 | 1 |
| 0x14 | EQ | 3 | 2 | 1 |
| 0x15 | ISZERO | 3 | 1 | 1 |
| 0x16 | AND | 3 | 2 | 1 |
| 0x17 | OR | 3 | 2 | 1 |
| 0x18 | XOR | 3 | 2 | 1 |
| 0x19 | NOT | 3 | 1 | 1 |
| 0x1a | BYTE | 3 | 2 | 1 |
| 0x20 | SHA3 | 30 | 2 | 1 |
| 0x30 | ADDRESS | 2 | 0 | 1 |
| 0x31 | BALANCE | 400 | 1 | 1 |
| 0x32 | ORIGIN | 2 | 0 | 1 |
| 0x33 | CALLER | 2 | 0 | 1 |
| 0x34 | CALLVALUE | 2 | 0 | 1 |
| 0x35 | CALLDATALOAD | 3 | 1 | 1 |
| 0x36 | CALLDATASIZE | 2 | 0 | 1 |
| 0x37 | CALLDATACOPY | 3 | 3 | 0 |
| 0x38 | CODESIZE | 2 | 0 | 1 |
| 0x39 | CODECOPY | 3 | 3 | 0 |
| 0x3a | GASPRICE | 2 | 0 | 1 |
| 0x3b | EXTCODESIZE | 700 | 1 | 1 |
| 0x3c | EXTCODECOPY | 700 | 4 | 0 |

| | | | | |
|------|----------------|-------|---|---|
| 0x3d | RETURNDATASIZE | 2 | 0 | 1 |
| 0x3e | RETURNDATACOPY | 3 | 3 | 0 |
| 0x40 | BLOCKHASH | 20 | 1 | 1 |
| 0x41 | COINBASE | 2 | 0 | 1 |
| 0x42 | TIMESTAMP | 2 | 0 | 1 |
| 0x43 | NUMBER | 2 | 0 | 1 |
| 0x44 | DIFFICULTY | 2 | 0 | 1 |
| 0x45 | GASLIMIT | 2 | 0 | 1 |
| 0x50 | POP | 2 | 1 | 0 |
| 0x51 | MLOAD | 3 | 1 | 1 |
| 0x52 | MSTORE | 3 | 2 | 0 |
| 0x53 | MSTORE8 | 3 | 2 | 0 |
| 0x54 | SLOAD | 200 | 1 | 1 |
| 0x55 | SSTORE | 0 | 2 | 0 |
| 0x56 | JUMP | 8 | 1 | 0 |
| 0x57 | JUMPI | 10 | 2 | 0 |
| 0x58 | PC | 2 | 0 | 1 |
| 0x59 | MSIZE | 2 | 0 | 1 |
| 0x5a | GAS | 2 | 0 | 1 |
| 0x5b | JUMPDEST | 1 | 0 | 0 |
| 0xa0 | LOG0 | 375 | 2 | 0 |
| 0xa1 | LOG1 | 750 | 3 | 0 |
| 0xa2 | LOG2 | 1125 | 4 | 0 |
| 0xa3 | LOG3 | 1500 | 5 | 0 |
| 0xa4 | LOG4 | 1875 | 6 | 0 |
| 0xf0 | CREATE | 32000 | 3 | 1 |
| 0xf1 | CALL | 700 | 7 | 1 |
| 0xf2 | CALLCODE | 700 | 7 | 1 |
| 0xf3 | RETURN | 0 | 2 | 0 |
| 0xf4 | DELEGATECALL | 700 | 6 | 1 |
| 0xf5 | CALLBLACKBOX | 40 | 7 | 1 |
| 0xfa | STATICCALL | 40 | 6 | 1 |
| 0xfd | REVERT | 0 | 2 | 0 |
| 0xff | SUICIDE | 5000 | 1 | 1 |

References

- [1] D. G. Wood, *Ethereum: A secure decentralised generalised transaction ledger*, <https://github.com/ethereum/yellowpaper>, 2017 (cit. on pp. 6, 9–11).
- [2] etherscan.io, *Charts: Ethereum blocksize history*. [Online]. Available: <https://etherscan.io/chart/blocksize>.
- [3] S. Levy, *Hackers: Heroes of the Computer Revolution*. O'Reilly, 2010.
- [4] BillWagner, *Serialization (c#)*. [Online]. Available: <https://docs.microsoft.com/en-us/dotnet/csharp/programming-guide/concepts/serialization/> (cit. on p. 17).
- [5] *Design patterns and refactoring*. [Online]. Available: https://sourcemaking.com/design_patterns/singleton (cit. on p. 17).
- [6] G. E. Ngondi and A. Butterfield, *A dictionary of computer science*. Oxford University Press, 2016 (cit. on p. 17).
- [7] E. Foundation, *Ethereum whitepaper*, <https://github.com/ethereum/wiki/wiki/White-Paper>, 2017 (cit. on pp. 6, 17).

Glossary

Address A 160-bit (20-byte) code used for identifying Accounts. 16

addresses 20 character strings, specifically the right-most 20 characters of the Keccak-256 hash of the RLP-derived mapping which contains the sender's address and the nonce of the block.. 16

balance A value which is intrinsic to accounts; the quantity of Wei in the account. All EVM operations are associated with changes in account balance. 16

beneficiary The 20-character (160-bit) address to which all fees collected from the successful mining of this block be transferred. 16

Bit The smallest unit of electronic data storage: there are eight bits in one byte. The Yellowpaper gives certain values in bits (e.g. 160 bits instead of 20 bytes). 16

block header Any information in a block besides transaction information. 16

Contract A piece of EVM Code that may be associated with an Account or an Autonomous Object. 16

Cryptographic hashing functions Hash functions make secure blockchains possible by establishing universal inputs for which there can only be one given output.^aThe reason this works is because the hash of a block's data is a certainty, just like two plus two equals four is a certainty.. 16

Ethereum Runtime Environment The environment which is provided to an Autonomous Object executing in the EVM. Includes the EVM but also the structure of the world state on which the relies for certain I/O instructions including CALL & CREATE. 16

Ethereum Foundation The non-profit organization in charge of executing the development processes of Ethereum in line with the [Whitepaper](#). 16

Ethereum Virtual Machine A sub-process of the *State Transition Function* which initializes and executes all of the transactions (ergo computations) in a block, prior to their finalization into the state.. 16

EVM Assembly The human readable version of EVM code. 16

EVM Code The bytecode that the EVM can natively execute. Used to formally specify the meaning and ramifications of a message to an Account. 16

Gas The fundamental network cost unit; gas is paid for exclusively by Ether. 16

leaf node the bottom-most node in a particular tree, of blocks, one half of the "key" the other half being the root node, which creates the path between. 16

Lower-Level Lisp The Lisp-like Low-level Language, a human-writable language used for authoring simple contracts and general low-level language toolkit for trans-compiling to. 16

Message Data (as a set of bytes) and Value (specified in Wei) that is passed between two accounts.. 16

Recursive Length Prefix Recursive Length Prefix. 16

root node the uppermost node in a particular tree, of blocks, representing a single world state^σ at a particular time. 5, 16

serialization Serialization is the process of converting an object into a stream of bytes in order to store the object or transmit it to memory, a database, or a file. Its main purpose is to save the state of an object in order to be able to recreate it when

^aActually, most hashing functions eventually have some collision points where two viable inputs reproduce the same output. But actual collision points are rare discoveries and tend to be followed (if not preceded by) newer more powerful hashing algorithms that are yet harder to break or find collisions in. Since the number space is infinite, we aren't likely to run out of potential new and larger hashing algorithms any time soon. Older hashing algorithms with known collisions, such as MD5 are not recommended for use in applications with stringent security requirements.

needed. The reverse process is called deserialization.[4] In Ethereum, most data is serialized through tree structures.. 16

singleton A design pattern in Object-Oriented Programming which specifies a class with one instance but with a global point of access to it[5]. 16

specification Technical descriptions, instructions, and definitions from which other people can create models. 16

state machine The term *State Machine* is reserved for any simple or complex process that moves deterministically from one discrete state to the next.. 16

state database A database stored off-chain, [i.e. on the computer of some user running an Ethereum client] which contains a trie structure mapping bytearrays [i.e. organized chunks of binary data] to other bytearrays [other organized chunks of binary data]. The *relationships* between each node on this trie constitute a MAP, a.k.a. a MAP-PING of all previous *world states* which a client might need to reference. 5, 7, 16

storage root One aspect of an ACCOUNT'S STATE: this is the hash of the trie^a that decides the STORAGE CONTENTS of the account. 16

Storage State The information particular to a given account that is maintained between the times that the account's associated EVM Code runs. 16

transaction A piece of data, signed by an External Actor. It represents either a Message or a new Autonomous Object. Transactions are recorded into each block of the blockchain. A transaction can also be an input message to a system that, because of the nature of the real-world event or activity it reflects, is required to be regarded as a single unit of work guaranteeing to either be processed completely or not at all.[6]. 16

Transaction . 16

trie A tree-structure for organizing data, the position of data in the tree contains the particular path

from root to leaf node that represents the key (the path from root to leaf is "one" key) you are searching the trie structure for. The data of the key is contained in the trie relationships that emerge from related nodes in the trie structure. 5, 7, 16

Whitepaper A conceptual map, distinct from the Yellowpaper, which highlights the development goals for Ethereum as a whole[7]. 16

Yellowpaper Ethereum's primary formal specification, written by Dr. Gavin Wood, one of the founders of Ethereum.. 16

Acronyms

ERE Ethereum Runtime Environment. 16

EVM Ethereum Virtual Machine. 16

LLL Lower Level Lisp. 16

^aA particular path from root to leaf in the state database